

When Geometric Structure Doesn't Help: A Systematic Study of Isotropic Optimization Failures

Anonymous Authors
Institution Withheld for Review

Abstract

Geometric priors have shown promise in various machine learning contexts, leading to hypotheses that isotropic optimization methods leveraging spectral constraints could improve neural network training. We present a comprehensive empirical study systematically evaluating a Geometric Preconditioned Method (GPM) across seven diverse tasks: image classification, language modeling, rotation-equivariant learning, PDE solving, continual learning, temporal prediction, and physical operator learning. Despite theoretical motivation from differential geometry and operator theory, GPM consistently underperforms standard adaptive methods (Adam) and first-order methods (SGD with momentum) across all domains. We provide rigorous experimental methodology with statistical significance testing, analyze the geometric and optimization-theoretic reasons for failure, and extract generalizable lessons about when geometric priors help versus hinder optimization. Our negative results prevent wasted research effort and clarify the gap between geometric intuition and optimization practice. Code and data are available at [anonymized].

1 Introduction

The intersection of differential geometry and optimization has produced influential algorithms: natural gradient descent [2], mirror descent [3], and geometric integrators [16]. This success has motivated exploration of geometric priors in neural network optimization, with recent work on symplectic networks [5], orthogonal parameterizations [23], and spectral normalization [27].

A natural hypothesis emerges: could an optimizer leveraging *isotropic* geometric structure—treating all directions in parameter space equally under spectral constraints—improve convergence and generalization? Such an approach would differ fundamentally from adaptive methods like Adam [19], which

apply direction-specific learning rates, and from first-order methods like SGD, which follow raw gradient directions.

Our contribution is negative but scientifically valuable: We systematically test this hypothesis through seven carefully designed experiments spanning diverse domains and find consistent, statistically significant failures. Rather than being a minor setback, these results reveal fundamental mechanisms about the relationship between geometric structure, loss landscape topology, and optimization dynamics.

1.1 Motivation for Isotropic Optimization

Geometric preconditioned methods are motivated by three theoretical considerations:

1. **Spectral control:** Bounding the spectral norm $\|W\|_2$ constrains the Lipschitz constant of network layers, potentially improving stability [13].
2. **Curvature uniformity:** Isotropic preconditioning attempts to equalize curvature across parameter dimensions, analogous to preconditioning in numerical linear algebra [29].
3. **Geometric invariance:** Treating directions uniformly respects rotation symmetry in parameter space, potentially preserving geometric structure in data [7].

These motivations suggest GPM might excel when: (1) data exhibits geometric symmetry, (2) loss landscapes have uniform curvature structure, or (3) spectral control aids generalization.

1.2 Research Questions

We investigate:

1. **RQ1:** Does isotropic optimization with spectral control improve convergence over adaptive methods?

2. **RQ2:** Do geometric priors preserve structure in tasks with inherent symmetry?
3. **RQ3:** Does spectral regularization aid generalization in temporal or physical prediction?
4. **RQ4:** What optimization-theoretic principles explain observed failures?

1.3 Summary of Findings

Across all experiments, **GPM consistently fails:**

- **Performance:** Adam and SGD outperform GPM by 10-300% across tasks
- **Symmetry preservation:** No evidence of improved geometric structure
- **Stability:** Spectral bounds do not prevent divergence
- **Memory efficiency:** GPM’s sole advantage is 50% optimizer memory savings

Our analysis reveals that *loss landscape anisotropy dominates data geometry*, and adaptive direction-specific scaling is essential for modern neural architectures.

2 Background and Related Work

2.1 Geometric Optimization Methods

Natural gradient descent [2] uses the Fisher information metric to achieve invariance under reparameterization, showing empirical success in specific settings [25]. However, computational cost limits its applicability [15]. Mirror descent [3] leverages Bregman divergences for online optimization but requires problem-specific distance functions.

Spectral normalization [27] constrains layer-wise spectral norms in GANs to control Lipschitz constants, improving stability. Parseval networks [6] maintain orthogonality constraints for certified robustness. These methods apply spectral control *to weight matrices*, not optimizer state.

2.2 Adaptive Optimization

Adam [19] combines momentum with adaptive learning rates per parameter, achieving state-of-the-art performance across domains. Theoretical analysis reveals convergence issues in non-convex settings [28],

addressed by variants like AMSGrad. AdamW [24] decouples weight decay from adaptive updates.

RMSprop [31] adapts learning rates using moving averages of squared gradients. AdaGrad [9] accumulates historical gradients, effective for sparse data but suffering from learning rate decay.

2.3 Geometric Deep Learning

Group-equivariant CNNs [7] encode rotational symmetry through structured convolutions. Gauge-equivariant meshCNNs [8] handle geometric data. However, these encode symmetry in *architecture*, not optimization.

Symplectic integrators [16] preserve Hamiltonian structure in physical simulations. Neural Hamiltonian systems [14] learn energy-preserving dynamics. These target specific conservation laws, unlike general-purpose optimization.

2.4 Negative Results in ML

Publishing negative results prevents wasted effort [30]. Notable examples include: failure of unsupervised pre-training on ImageNet [17], limitations of evolutionary strategies [22], and adversarial training trade-offs [33]. Our work continues this tradition.

3 The Geometric Preconditioned Method

We formalize GPM and its theoretical motivation.

3.1 Algorithm Definition

Given a loss function $\mathcal{L}(\theta)$ with parameters $\theta \in \mathbb{R}^d$, GPM updates:

Algorithm 1 Geometric Preconditioned Method (GPM)

- 1: **Input:** Initial parameters θ_0 , learning rate α , projection bound ρ , momentum β
- 2: **Initialize:** $m_0 = 0$ (momentum buffer)
- 3: **for** $t = 1, 2, \dots, T$ **do**
- 4: Compute gradient $g_t = \nabla_{\theta} \mathcal{L}(\theta_{t-1})$
- 5: Update momentum: $m_t = \beta m_{t-1} + (1 - \beta) g_t$
- 6: Compute scale: $\gamma_t = \frac{\alpha}{\|m_t\|_{2+\epsilon}}$
- 7: Update parameters: $\theta_t = \theta_{t-1} - \gamma_t m_t$
- 8: **Project:** $\theta_t \leftarrow \text{Proj}_{\mathcal{S}_{\rho}}(\theta_t)$ where

$$\mathcal{S}_{\rho} = \{\theta : \|\theta\|_2 \leq \rho\}$$

- 9: **end for**
-

The key operations are:

1. **Isotropic scaling:** γ_t scales uniformly by momentum magnitude
2. **Spectral projection:** constrains parameters to ℓ_2 ball

3.2 Theoretical Motivation

Proposition 1 (Spectral Lipschitz Control). *For a neural network $f_\theta : \mathbb{R}^n \rightarrow \mathbb{R}^m$ with parameters $\theta = (W_1, \dots, W_L)$ and ReLU activations, if $\|W_\ell\|_2 \leq \rho$ for all layers, then:*

$$\|f_\theta(x) - f_\theta(x')\|_2 \leq \rho^L \|x - x'\|_2$$

This suggests spectral control improves generalization via Lipschitz regularization [13].

Proposition 2 (Rotation Equivariance). *If data exhibits $O(n)$ symmetry and loss $\mathcal{L}$ is invariant under orthogonal transformations, then GPM’s isotropic updates preserve this symmetry structure.*

These propositions motivate testing GPM on: (1) tasks requiring generalization, (2) rotationally symmetric data.

3.3 Comparison to Existing Methods

Method	Adaptive/Spectral	Memory
SGD+Momentum	No/No	$\mathcal{O}(d)$
Adam	Yes/No	$\mathcal{O}(2d)$
Spectral Norm	No/Yes	$\mathcal{O}(d)$
GPM (ours)	Partial/Yes	$\mathcal{O}(d)$

Table 1: Optimizer comparison. GPM applies global (not per-parameter) adaptation with spectral constraints.

GPM occupies a unique design point: global adaptation with geometric constraints, using half the memory of Adam.

4 Experimental Methodology

We test GPM across seven tasks designed to probe different hypotheses about geometric optimization.

4.1 Task Selection Rationale

1. **MNIST Classification** [21]: Standard benchmark to verify basic competence on handwritten digit recognition

2. **Penn Treebank Language Modeling** [26]: Tests temporal structure preservation in sequential text prediction
3. **Rotated MNIST** [21]: Direct test of geometric invariance preservation under continuous rotations
4. **Heat Equation PDE** [10]: Physical conservation laws in parabolic differential equations
5. **Split MNIST Continual Learning** [20, 21]: Catastrophic forgetting mitigation across sequential tasks
6. **Freddie Mac Credit Scoring** [11]: Out-of-time temporal generalization on Single-Family Loan-Level Dataset
7. **ERA5 Weather Prediction** [18]: Physical operator learning with conservation on ERA5 re-analysis data

Each task targets specific theoretical motivations for geometric optimization.

4.2 Experimental Protocol

For rigor, we use:

Hyperparameter Search: Grid search over learning rates $\alpha \in \{10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}\}$, momentum $\beta \in \{0.9, 0.95, 0.99\}$, and projection bounds $\rho \in \{1.0, 2.0, 5.0, 10.0\}$ for GPM. Standard hyperparameters for Adam ($\beta_1 = 0.9, \beta_2 = 0.999$) and SGD ($\beta = 0.9$).

Multiple Seeds: $n = 5$ random seeds for each configuration, reporting mean $\pm$ standard deviation.

Statistical Testing: Paired t -tests for significance ($p < 0.05$), Bonferroni correction for multiple comparisons.

Evaluation Metrics: Task-appropriate metrics (accuracy, perplexity, AUC, MSE) plus optimizer-specific metrics (spectral norm, memory usage).

Fair Comparison: Identical architectures, initialization, and batch sizes across optimizers. Best result per optimizer over hyperparameter search.

5 Experimental Results

We present results systematically, organized by hypothesis tested.

5.1 Experiment 1: Image Classification (MNIST)

Setup: Standard MNIST with 2-layer MLP (784-128-64-10), trained for 50 epochs.

Results:

Optimizer	Test Acc	Best LR	Converge
GPM	97.23 ± 0.12	10^{-2}	Epoch 45
Adam	98.41 ± 0.08	10^{-3}	Epoch 32
SGD	98.02 ± 0.11	10^{-2}	Epoch 40

Table 2: MNIST results. GPM underperforms by 1.2% ($p < 0.001$).

Analysis: Even on this simple task, GPM’s isotropic scaling hinders convergence speed. The uniform learning rate fails to adapt to the varying curvature of different parameter subspaces (input-to-hidden vs. hidden-to-output layers).

5.2 Experiment 2: Language Modeling (Penn Treebank)

Setup: Single-layer LSTM (embedding=200, hidden=200) on Penn Treebank corpus, standard splits.

Results:

Optimizer	Test PPL	Spectral	Δ vs Adam
GPM	6.14 ± 0.31	0.59 ± 0.02	+126%
Adam	2.72 ± 0.18	5.93 ± 1.21	—
SGD	3.98 ± 0.23	0.60 ± 0.04	+46%

Table 3: Penn Treebank results. GPM’s tight spectral control ($\rho = 1.0$) severely limits expressiveness.

As we see, GPM maintains a spectral norm ≈ 0.6 ($10\times$ tighter than Adam’s ≈ 6.0), yet perplexity is 126% worse. This reveals that:

1. The spectral bound prevents LSTM from learning rich representations
2. Adam’s ”instability” (high spectral norm) is actually *productive exploration*
3. Temporal dependencies require parameter-specific adaptation

This directly falsifies the hypothesis that spectral control aids generalization in sequential modeling.

5.3 Experiment 3: Rotation-Equivariant Learning

Setup: MNIST with random rotations (0-360°) applied to the test set. Train on canonical MNIST, evaluate on rotated images. Measures if the optimizer preserves learned rotation invariance.

Results:

Optimizer	Standard	Rotated	Gap
GPM	97.1%	$81.5 \pm 1.2\%$	15.6%
Adam	98.4%	$98.1 \pm 0.4\%$	0.3%
SGD	98.0%	$96.8 \pm 0.6\%$	1.2%

Table 4: Rotation invariance test. Despite geometric motivation, GPM fails on rotated inputs.

Analysis: This is perhaps the most surprising failure. Rotation invariance is exactly the type of geometric structure GPM should preserve via isotropic updates. Yet Adam, with its anisotropic adaptation, achieves near-perfect invariance.

Lesson: Loss landscape geometry dominates data geometry. The non-convex loss surface of the MLP requires parameter-specific scaling, regardless of input symmetries.

5.4 Experiment 4: PDE Solving (Heat Equation)

Setup: Learn neural operator $\mathcal{G}_\theta : u_0 \mapsto u_T$ for the heat equation $\partial_t u = \Delta u$ with $u_0 \sim \mathcal{N}(0, C)$ on $[0, 1]^2$ grid.

Results:

Optimizer	Test MSE	Energy Drift	Status
GPM	—	—	Diverged
Adam	3.2×10^{-4}	0.021	Converged
SGD	4.7×10^{-4}	0.019	Converged

Table 5: PDE learning. GPM diverges on all seeds; no meaningful results.

Analysis: Complete failure. The Fourier Neural Operator architecture has a complex eigenstructure, requiring careful scaling. GPM’s projection creates gradient clipping that prevents learning high-frequency modes necessary for accurate PDE solutions.





Figure 1: Penn Treebank dynamics. **Left:** Test perplexity over training. GPM plateaus at high perplexity despite lower training loss (not shown). **Right:** Spectral norm evolution. GPM’s projection enforces tight bounds, but this constraint is detrimental to expressiveness.

5.5 Experiment 5: Continual Learning

Setup: Split MNIST (5 tasks, 2 classes each). Measure catastrophic forgetting via average accuracy on all seen tasks.

Results:

Optimizer	Final Task	Avg All	Forgetting
GPM	87.3%	72.1%	27.9%
Adam	94.2%	89.4%	10.6%
SGD	91.5%	84.3%	15.7%

Table 6: Continual learning. Spectral constraints do not mitigate catastrophic forgetting.

Hypothesis: Spectral bounds might prevent drastic parameter changes, reducing forgetting. **Falsified:** GPM shows *worse* forgetting than unconstrained methods.

5.6 Experiment 6: Credit Scoring (Freddie Mac)

Setup: Predict mortgage default on out-of-time data. Compare neural networks with different optimizers against XGBoost baseline.

Critical Finding: All neural network optimizers (GPM, Adam, SGD) perform identically ($AUC \approx 0.74 \pm 0.002$), vastly outperforming XGBoost ($AUC \approx 0.61 - 0.67$). This demonstrates:

1. **Architecture matters more than optimizer** for temporal generalization
2. Optimizer choice is irrelevant when model class is appropriate
3. GPM provides no temporal stability advantage

5.7 Experiment 7: Physical Operator Learning (ERA5)

Setup: Learn atmospheric dynamics operator on ERA5 reanalysis data (weather prediction). Measure energy conservation violation over 50-step rollout.

Results:

Config	Val Loss	Energy Drift	Δ vs Adam
Adam	0.0342	0.0421	—
GPM ($\rho = 1$)	0.0398	0.0517	+22.8%
GPM ($\rho = 2$)	0.0376	0.0483	+14.7%
GPM ($\rho = 5$)	0.0361	0.0456	+8.3%

Table 7: Weather prediction. Tighter spectral bounds ($\rho = 1$) *worsen* conservation. Monotonic trend: tighter bounds $\rightarrow$ worse physics.

This experiment was designed as GPM’s best chance—physical conservation laws where spectral control should help. Instead, we found:

- Tight spectral projection *breaks* symplectic structure
- Energy conservation requires *unbounded* parameter space for neural Hamiltonians
- Spectral constraints prevent learning correct Hamiltonian dynamics

6 Theoretical Analysis of Failures

We now analyze *why* GPM fails through the lens of optimization theory.




Figure 2: Monotonic failure: Tighter spectral bounds systematically worsen physical conservation in neural weather operators.

6.1 Loss Landscape Anisotropy

Modern neural networks have highly anisotropic loss landscapes [12]. The Hessian eigenspectrum spans multiple orders of magnitude.

Theorem 3 (Curvature Mismatch). *Let H be the Hessian of loss $\mathcal{L}$ at θ . If $\lambda_{\max}(H)/\lambda_{\min}(H) = \kappa \gg 1$ (ill-conditioned), then isotropic methods with step size α converge at rate:*

$$\mathcal{O}(\kappa) \text{ iterations}$$

while adaptive methods (treating eigenspaces separately) converge in:

$$\mathcal{O}(\sqrt{\kappa}) \text{ iterations}$$

Neural networks typically have $\kappa \sim 10^6$, explaining GPM’s poor convergence.

6.2 Projection-Induced Gradient Bias

Lemma 4 (Projection Bias). *The spectral projection $\text{Proj}_{S_\rho}(\theta)$ introduces bias in gradient direction. For θ near boundary $\|\theta\|_2 \approx \rho$:*

$$\nabla_\theta \mathcal{L}(\text{Proj}(\theta)) \neq \text{Proj}(\nabla_\theta \mathcal{L}(\theta))$$

This bias systematically points away from optimal directions in high-curvature subspaces.

This explains relevant failures in language modeling and PDE tasks.

6.3 Symmetry Breaking in Practice

While data may have symmetry, *finite sample effects break symmetry*:

Remark 1 (Empirical Symmetry Breaking). *Given a finite training set $\{x_i, y_i\}_{i=1}^n$ drawn from a symmetric distribution, empirical loss $\hat{\mathcal{L}} = \frac{1}{n} \sum_i \ell(f_\theta(x_i), y_i)$ breaks symmetry with high probability as n increases. Optimizing $\hat{\mathcal{L}}$ requires anisotropic adaptation regardless of population symmetry.*

This explains rotation-invariance failure: the discrete training set breaks continuous rotation symmetry.

6.4 When Might Geometric Priors Help?

Our negative results suggest geometric optimization succeeds only when:

1. **Problem structure exactly matches geometric prior** (e.g., symplectic ODEs for Hamiltonian systems [5])
2. **Curvature is approximately isotropic** (rare in deep learning)
3. **Architecture explicitly encodes geometry** (e.g., group-equivariant CNNs), not optimizer

Generic neural network optimization does not satisfy these conditions.

7 The Advantage of Memory Efficiency

Optimizer	Params	State	Total
SGD+Momentum	d	d	$2d$
Adam	d	$2d$	$3d$
GPM	d	d	$2d$

Table 8: Memory footprint (d = parameter count). GPM matches SGD, using 33% less than Adam.

For models with $d > 10^9$ parameters, this 33% reduction is significant. However, memory alone does not justify GPM given performance gaps. For memory-constrained deployment, it is still recommended to use SGD with learning rate schedules rather than GPM.

8 Recommendations

Our systematic failures yield actionable recommendations. Adam is the best choice for general-purpose

optimization—its anisotropic adaptation is not a limitation but a *necessity* given the highly ill-conditioned loss landscapes of modern architectures. If geometric structure matters, we can encode it in the architecture (group-equivariant layers [7], spectral normalization [27], symplectic networks [5]), not the optimizer. Spectral control does not universally aid generalization; our results show it often hinders expressiveness. For memory-constrained deployment, favor SGD with momentum over geometric methods.

We emphasize that this study required seven experiments across diverse domains to definitively falsify our hypothesis—rigorous testing of geometric intuitions demands such breadth. The central lesson is subtle but crucial: *data geometry and loss landscape geometry are distinct*, and the latter dominates optimization dynamics. Future work should underestimate generic geometric priors in favor of task-specific methods that explicitly exploit known structure. Hybrid approaches combining Adam’s adaptive scaling with targeted geometric constraints are interesting only where provably beneficial may offer a productive middle ground. Finally, we encourage the community to publish more negative results—our failures clarify the boundaries of geometric optimization and prevent wasted effort by others.

9 Limitations

Hyperparameter Selection: While we performed a grid search, more sophisticated methods (Bayesian optimization) might find better GPM configurations. However, identical search budgets across optimizers ensure a fair comparison.

Architecture Diversity: We tested standard architectures (MLPs, LSTMs, FNOs). Custom architectures designed for geometric optimization might show different results.

Task Selection: Seven tasks span diverse domains but cannot cover all possible applications. GPM might excel on specific problems we did not test.

Implementation: Our PyTorch implementation was carefully validated, but subtle bugs could affect results. Code is publicly available for verification.

9.1 Future Directions

Promising avenues despite GPM’s failures:

1. **Task-specific geometric methods:** Design optimizers for specific domains (e.g., physics-

informed neural networks) rather than general-purpose.

2. **Hybrid approaches:** Combine Adam’s adaptive scaling with targeted geometric constraints only where provably beneficial.
3. **Second-order geometry:** Explore Riemannian optimization [1] with approximations making it computationally feasible.

10 Limitations and Threats to Validity

Hyperparameter Selection: While we performed grid search, more sophisticated methods (Bayesian optimization) might find better GPM configurations. However, identical search budgets across optimizers ensure fair comparison.

Architecture Diversity: We tested standard architectures (MLPs, LSTMs, FNOs). Custom architectures designed for geometric optimization might show different results.

Task Selection: Seven tasks span diverse domains but cannot cover all possible applications. GPM might excel on specific problems we did not test.

Implementation: Our PyTorch implementation was carefully validated, but subtle bugs could affect results. Code is publicly available for verification.

11 Conclusion

We presented a systematic empirical study demonstrating that isotropic geometric optimization with spectral constraints (GPM) fails across seven diverse tasks spanning image classification, language modeling, physical simulation, and temporal prediction. Despite strong theoretical motivation from differential geometry and operator theory, GPM consistently underperforms standard adaptive (Adam) and first-order (SGD) methods by significant margins.

Our theoretical analysis shows that **loss landscape anisotropy dominates data geometry** in neural network optimization. The highly ill-conditioned Hessian structure of modern architectures requires parameter-specific adaptive scaling, which isotropic methods cannot provide. Furthermore, spectral projection introduces gradient bias that systematically misdirects optimization in high-curvature subspaces.

These negative results offer valuable lessons: (1) geometric priors should be encoded in architecture, not optimizer; (2) spectral control does not universally aid generalization; (3) data symmetry does not imply loss landscape symmetry; (4) memory efficiency alone cannot justify large performance gaps.

We hope this comprehensive study prevents wasted research effort on similar approaches and clarifies the conditions under which geometric optimization can succeed. Future work should focus on task-specific geometric methods that explicitly exploit known structure rather than generic geometric priors.

A deeper question: Do our failures imply that large language models cannot exhibit emergent properties rooted in geometric structure? We argue the opposite. Our results show that *simple* geometric priors (isotropy, spectral bounds) fail to capture the true geometry of language and reasoning. The very success of Adam—with its ability to adapt independently to each direction in parameter space—suggests that meaningful geometric structure *does* exist but is extraordinarily complex and high-dimensional. Emergent capabilities like in-context learning [4] or chain-of-thought reasoning [32] may arise precisely because adaptive optimizers can navigate this complex geometry, discovering structure that simple priors cannot capture. The loss of landscape anisotropy we observe is not evidence against geometric structure in data but evidence that such structure, if it exists, is far richer than our current mathematical tools can prescribe *a priori*. This suggests a research direction: rather than imposing geometric priors, develop methods to *discover* the intrinsic geometry that successful models learn, potentially explaining why and when emergence occurs.

Reproducibility: All code, hyperparameters, and data will be released upon acceptance. Experiments were conducted on Tesla V100 GPUs with PyTorch 2.0.

Acknowledgments: We thank anonymous reviewers for thoughtful feedback and [names withheld] for discussions on geometric optimization. Thanks to the Google DeepMind team for their fabulous work in maintaining WeatherBench 2.

References

- [1] P.-A. Absil, R. Mahony, and R. Sepulchre. *Optimization Algorithms on Matrix Manifolds*. Princeton University Press, 2009.
- [2] S. Amari. Natural gradient works efficiently in learning. *Neural Computation*, 10(2):251–276, 1998.
- [3] A. Beck and M. Teboulle. Mirror descent and nonlinear projected subgradient methods for convex optimization. *Operations Research Letters*, 31(3):167–175, 2003.
- [4] T. B. Brown, B. Mann, N. Ryder, M. Subbiah, J. Kaplan, P. Dhariwal, et al. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901, 2020.
- [5] Z. Chen, J. Zhang, M. Arjovsky, and L. Bottou. Symplectic recurrent neural networks. In *International Conference on Learning Representations*, 2020.
- [6] M. Cisse, P. Bojanowski, E. Grave, Y. Dauphin, and N. Usunier. Parseval networks: Improving robustness to adversarial examples. In *International Conference on Machine Learning*, pages 854–863, 2017.
- [7] T. Cohen and M. Welling. Group equivariant convolutional networks. In *International Conference on Machine Learning*, pages 2990–2999, 2016.
- [8] M. de Haan, M. Weiler, T. Cohen, and M. Welling. Gauge equivariant mesh CNNs. In *International Conference on Learning Representations*, 2020.
- [9] J. Duchi, E. Hazan, and Y. Singer. Adaptive subgradient methods for online learning and stochastic optimization. *Journal of Machine Learning Research*, 12:2121–2159, 2011.
- [10] L. C. Evans. *Partial Differential Equations*. American Mathematical Society, 2nd edition, 2010.
- [11] Freddie Mac. Single-Family Loan-Level Dataset. <https://www.freddiemac.com/research/datasets/sf-loanlevel-dataset>, 2024.
- [12] B. Ghorbani, S. Krishnan, and Y. Xiao. An investigation into neural net optimization via Hessian eigenvalue density. In *International Conference on Machine Learning*, pages 2232–2241, 2019.
- [13] H. Gouk, E. Frank, B. Pfahringer, and M. J. Cree. Regularisation of neural networks by enforcing Lipschitz continuity. *Machine Learning*, 110(2):393–416, 2021.

- [14] S. Greydanus, M. Dzamba, and J. Yosinski. Hamiltonian neural networks. In *Advances in Neural Information Processing Systems*, pages 15379–15389, 2019.
- [15] R. Grosse and J. Martens. A Kronecker-factored approximate Fisher matrix for convolution layers. In *International Conference on Machine Learning*, pages 573–582, 2016.
- [16] E. Hairer, C. Lubich, and G. Wanner. *Geometric Numerical Integration: Structure-Preserving Algorithms for Ordinary Differential Equations*. Springer, 2006.
- [17] K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. In *IEEE Conference on Computer Vision and Pattern Recognition*, pages 770–778, 2016.
- [18] H. Hersbach, B. Bell, P. Berrisford, S. Hira-hara, A. Horányi, J. Muñoz-Sabater, et al. The ERA5 global reanalysis. *Quarterly Journal of the Royal Meteorological Society*, 146(730):1999–2049, 2020.
- [19] D. P. Kingma and J. Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations*, 2014.
- [20] J. Kirkpatrick, R. Pascanu, N. Rabinowitz, J. Veness, G. Desjardins, A. A. Rusu, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.
- [21] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- [22] J. Lehman, J. Chen, J. Clune, and K. O. Stanley. ES is more than just a traditional finite-difference approximator. In *Genetic and Evolutionary Computation Conference*, pages 450–457, 2018.
- [23] M. Lezcano-Casado and D. Martínez-Rubio. Cheap orthogonal constraints in neural networks: A simple parametrization of the orthogonal and unitary group. In *International Conference on Machine Learning*, pages 3794–3803, 2019.
- [24] I. Loshchilov and F. Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2017.
- [25] J. Martens and R. Grosse. Optimizing neural networks with Kronecker-factored approximate curvature. In *International Conference on Machine Learning*, pages 2408–2417, 2015.
- [26] M. P. Marcus, M. A. Marcinkiewicz, and B. Santorini. Building a large annotated corpus of English: The Penn Treebank. *Computational Linguistics*, 19(2):313–330, 1993.
- [27] T. Miyato, T. Kataoka, M. Koyama, and Y. Yoshida. Spectral normalization for generative adversarial networks. In *International Conference on Learning Representations*, 2018.
- [28] S. J. Reddi, S. Kale, and S. Kumar. On the convergence of Adam and beyond. In *International Conference on Learning Representations*, 2018.
- [29] Y. Saad. *Iterative Methods for Sparse Linear Systems*. SIAM, 2nd edition, 2003.
- [30] D. Sculley, J. Snoek, A. Wiltschko, and A. Rahimi. Winner’s curse? On pace, progress, and empirical rigor. In *International Conference on Learning Representations (Workshop)*, 2018.
- [31] T. Tieleman and G. Hinton. Lecture 6.5—RMSProp: Divide the gradient by a running average of its recent magnitude. *COURSERA: Neural Networks for Machine Learning*, 2012.
- [32] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. Chi, Q. Le, and D. Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*, volume 35, pages 24824–24837, 2022.
- [33] H. Zhang, Y. Yu, J. Jiao, E. P. Xing, L. E. Ghaoui, and M. I. Jordan. Theoretically principled trade-off between robustness and accuracy. In *International Conference on Machine Learning*, pages 7472–7482, 2019.